

# Learn to Play CartPole with PyTorch using PGPE

Code: [agent.py](#), [cartpole.py](#)

See the [final result](#) on my OpenAI Gym profile

See the [sample code](#) on my Github

## INTRODUCTION

In this post, we will apply another policy-based algorithm to solve the CartPole environment. Unlike REINFORCE, searching the optimal policy in the action space, policy gradient with parameter-based explorations (PGPE) searches the optimal policy in the parameter space. The idea is to equip each model parameter with a *prior distribution* to sample from. At the beginning of each roll-out, we sample a set of model parameters, and follow a *deterministic* controller (policy) for the episode thereafter. Just like REINFORCE, the original PGPE adopts an episodic setting. After an episode finishes, we compute the gradient of the objective function w.r.t. *hyper-parameters of the prior distribution* and update, again, the hyper-parameters.

## ALGORITHM

*Note:* Some notations have been modified to better fit in the context, hence you may find some inconsistency in notations with the original paper.

1. Each action follows the *deterministic policy* (the  $\delta$  term) with model parameters  $\delta$  drawn from the *prior distribution*, parametrized by  $\rho$

$$p(a|s, \rho) = \int_{\Theta} p(\theta|\rho) \delta_{F_{\theta}(s)=a} d\theta$$

which means the probability of taking action  $a$  under state  $s$  is exactly the probability of drawing parameter  $\theta$  given hyper-parameter  $\rho$ , since the delta function (our deterministic controller) contributes no stochasticity in the equation.

2. The objective  $J(\rho)$  is defined to be

$$J(\rho) = \int_{\Theta} \int_H p(h, \theta|\rho) r(h) dh d\theta$$

which is the *expected return* taken over all possible parameters we can draw from the prior and the history generated by following policy  $p(h, \theta)$ .

3. Apply the classic log-likelihood trick, we obtain the gradient of the objective w.r.t. hyper-parameter  $\rho$

$$\nabla_{\rho} J(\rho) = \int_{\Theta} \int_H p(h, \theta|\rho) \nabla_{\rho} \log p(h, \theta|\rho) r(h) dh d\theta$$

4. Observe that the history is conditionally independent of  $\rho$  given  $\theta$ , that is,  $p(h, \theta | \rho) = p(h | \rho)p(\theta | \rho)$  and the partial derivative reduces,  $\nabla_{\rho} \log p(h, \theta | \rho) = \nabla_{\rho} \log p(\theta | \rho)$ . Combining all of the above, the gradient estimator reveals:

$$\nabla_{\rho} J(\rho) \approx \frac{1}{N} \sum_{n=1}^N \nabla_{\rho} \log p(\theta^n | \rho) r(h^n)$$

**Remark** the above gradient estimator is completely determined by the agent's model parameters *without* the knowledge of environment dynamics. Information concerning environment dynamics encapsulates in the term  $p(h | \theta)$ , as we get rid of it when taking the log value of the term and differentiate w.r.t.  $\rho$ . Hence PGPE is a *model-free* algorithm, just as REINFORCE.

5. Here we assume each model parameter  $\theta_i$  follows a normal distribution independent with one another. Hence the hyper-parameters  $\rho$  consists of a sequence of means and standard deviations  $\rho = ((\mu_i, \sigma_i)_i)$ . Direct computation gives

$$\nabla_{\mu} \log p(\theta | \rho) = \frac{\theta - \mu}{\sigma^2}, \quad \nabla_{\sigma} \log p(\theta | \rho) = \frac{(\theta - \mu)^2 - \sigma^2}{\sigma^3}$$

Finally, we choose a stepsize proportional to the variance,  $\alpha_{\rho} = \alpha \sigma^2$ . The stepsize may be different for different parameters as we will see below.

$$\Delta \mu = \alpha_{\mu} (r - b)(\theta - \mu), \quad \Delta \sigma = \alpha_{\sigma} (r - b) \frac{(\theta - \mu)^2 - \sigma^2}{\sigma}$$

where  $b$  is the reward baseline for variance reduction purpose.

## CODE

For brevity, I only explain the code snippet I think may raise difficulty to the reader and omit most part of it. Full code can be found in my GitHub.

***cartpole.py***

```

1  #...
2
3  # Train for 200 episodes
4  for ep in range(200):
5      observation = env.reset()
6
7      while True:
8          # The typical RL env-agent paradigm
9          action = RL.choose_action(observation)
10         observation_, reward, done, info = env.step(action)
11
12         RL.store_reward(reward)
13
14         if done:
15             # Learn in an episodic basis
16             vt = RL.learn_and_sample()
17             break
18
19         observation = observation_
20
21 # Close env.
22 env.close()

```

### **agent.py**

```

1  #...
2
3  class PGPE:
4      def __init__(self, n_actions, n_features):
5          #...
6
7          # Prepare hyper-params and initialize model params
8          self.Param = list(self.model.parameters())
9          self.Mu = []
10         self.Sigma = []
11         for p in self.Param:
12             # initialize hyper-params
13             self.Mu.append(torch.normal(torch.zeros(p.size()),
14 torch.ones(p.size()))))
15             self.Sigma.append(2 * torch.ones(p.size()))
16
17         # Sample initial model params
18         p.data = torch.normal(self.Mu[-1], self.Sigma[-1])
19
20     def choose_action(self, obs):
21         # Scale input to (-1, 1):
22         # 1. if range is finite -> divide by range
23         # 2. if range is infinite -> take tanh
24         obs[0] /= 2.4

```

```

24         obs[1] = np.tanh(obs[1])
25         obs[2] /= 41.8
26         obs[3] = np.tanh(obs[3])
27
28         #...
29
30         # Deterministic policy, pick action with greater value
31         action = a[0].argmax()
32         return action
33
34     #...
35
36     def learn_and_sample(self):
37         # Scale return to [0, 1]
38         _r = self.ret / 200
39
40         # reset return tracker
41         self.ret = 0.
42
43         for i in range(len(self.Param)):
44             # Learning
45             # These are the T and S matrices in the original paper
46             _T = self.Param[i].data - self.Mu[i]
47             _S = (_T ** 2 - self.Sigma[i] ** 2) / self.Sigma[i]
48
49             # Update means
50             _delta_Mu = self.Mu_lr * _r * _T
51             self.Mu[i] += _delta_Mu
52
53             # Update standard deviations
54             _delta_Sigma = self.Sigma_lr * _r * _S
55             self.Sigma[i] += _delta_Sigma
56
57             # Freeze params if hit target reward, else re-sample
58             if _r < 1.:
59                 self.Param[i].data = torch.normal(self.Mu[i],
self.Sigma[i])

```


## RESULT

The result of this algorithm can be viewed on [my OpenAI Gym page](#). I also provide a gist link there so you can easily reproduce the result. As you can see, the algorithm converges quickly, taking only 5 episodes to solve the CartPole environment. The model consists of only 8 parameters and without performing any sophisticated initialization trick (in fact, random initialization from  $\text{normal}(0, 2)$  is used to draw initial model parameters). Other results provided on the [OpenAI Gym cartpole environment page](#) utilizing the REINFORCE algo often take up dozens more episodes to converge to a stable policy.

By the way, in case anyone is interested in what the input states are in the cartpole env. Here is the [official OpenAI gym github page](#) providing the background information. Note the env setting is a bit different from the one on the wiki page. But it doesn't really matter since we scale the input state when it comes in.

## ROBUSTNESS

To further evaluate our algorithm, we roll out 30 trials, each with 200 episodes. Then compute the mean and standard deviation of reward **at each episode** over these 30 trials to get a sense of how much reward in average we can get for each episode. The plot demonstrates the result:



To interpret, say, at the 50th episode, the algorithm gets an average reward of 187.5 and reward standard deviation being 37.5. The plot demonstrates the desired outcome: **mean reward grows steadily toward 200 and the standard deviation decreases to 0 as the algorithm learns to prevent the pole from falling for 200 steps.** Also, we can see the algorithm *solves* the env. at around 63th step, where it achieves an average reward of 195.0 over 100 consecutive episodes thereafter.